

# Cluster API: The Future of Kubernetes Installers

Fabrizio Pandini

VMware  
Member of Technical Staff  
Italy



kubernetes

# Agenda


- Introducing Cluster API
  - How it works?
  - Demo



# Trivia Time!

## History of kubernetes

1. What was the first Kubernetes installer?
2. Name a Kubernetes installer or distribution?



# Cluster API



The Cluster API is a Kubernetes project to bring declarative, Kubernetes-style APIs to cluster creation, configuration, and management.

<https://cluster-api.sigs.k8s.io/>




kubernetes

# 1. Open Source



<https://github.com/kubernetes-sigs/cluster-api>

## 2. Awesome logo



### 3. Community traction



[sigs.k8s.io/cluster-api/#provider-implementations](https://sigs.k8s.io/cluster-api/#provider-implementations)



kubernetes

## 4. Declarative



```
kubectl apply -f cluster.yaml
```



kubernetes



## 5. Same experience on any provider



```
kubectl apply -f cluster-on-aws.yaml
```

```
kubectl apply -f cluster-on-vsphere.yaml
```

```
kubectl apply -f cluster-on-azure.yaml
```

```
kubectl apply -f cluster-on-google.yaml
```

...



kubernetes

# Agenda

- Intro
- **How it works?**
- Demo



# Guess what...



If you know a little about Pod and Deployments in Kubernetes, then you already know how Cluster API works.



kubernetes

# API objects in Kubernetes



In Kubernetes API objects are **abstractions** of real objects e.g.

*A Pod object is an abstraction of a workload running in a container*

# API objects in Cluster API



Then, why not to define  
a **Machine** object  
as an abstraction of any  
infrastructure running in a  
Kubernetes Node?



kubernetes

# API objects in Cluster API



Cluster



Machine



MachineSet



MachineDeployment



MachineClass



Pod



ReplicaSet



Deployment



StorageClass



kubernetes

# In practice: Cluster CRDs



## my-cluster spec

```
apiVersion: cluster.x-k8s.io/v1alpha2
kind: Cluster
metadata:
  name: my-cluster
spec:
  clusterNetwork:
    pods:
      cidrBlocks: ["192.168.0.0/16"]

infrastructureRef:
  apiVersion: infrastructure.cluster.x-k8s.io/v1alpha2
  kind: AWSCluster
  name: my-cluster-aws
```

## my-cluster spec for AWS

```
apiVersion: infrastructure.cluster.x-k8s.io/v1alpha2
kind: AWSCluster
metadata:
  name: my-cluster-aws
spec:
  region: us-east-1
  sshKeyName: default
```



# In practice: Machine CRDs



## my-machine spec

```
apiVersion: cluster.x-k8s.io/v1alpha2
kind: Machine
metadata:
  name: my-machine
spec:
  version: v1.15.3

infrastructureRef:
  apiVersion: infrastructure.cluster.x-k8s.io/v1alpha2
  kind: AWSMachine
  name: my-machine-aws
```

## my-machine spec for AWS

```
apiVersion: infrastructure.cluster.x-k8s.io/v1alpha2
kind: AWSMachine
metadata:
  name: my-machine-aws
spec:
  instanceType: t3.large
  sshKeyName: ...
  ...
```



kubernetes



# In practice: Machine CRDs



## my-machine spec

```
apiVersion: cluster.x-k8s.io/v1alpha2
kind: Machine
metadata:
  name: my-machine
spec:
  version: v1.15.3
  infrastructureRef:
    ...
  bootstrap:
    configRef:
      apiVersion: bootstrap.cluster.x-k8s.io/v1alpha2
      kind: KubeadmConfig
      name: my-machine-bootstrapper
```

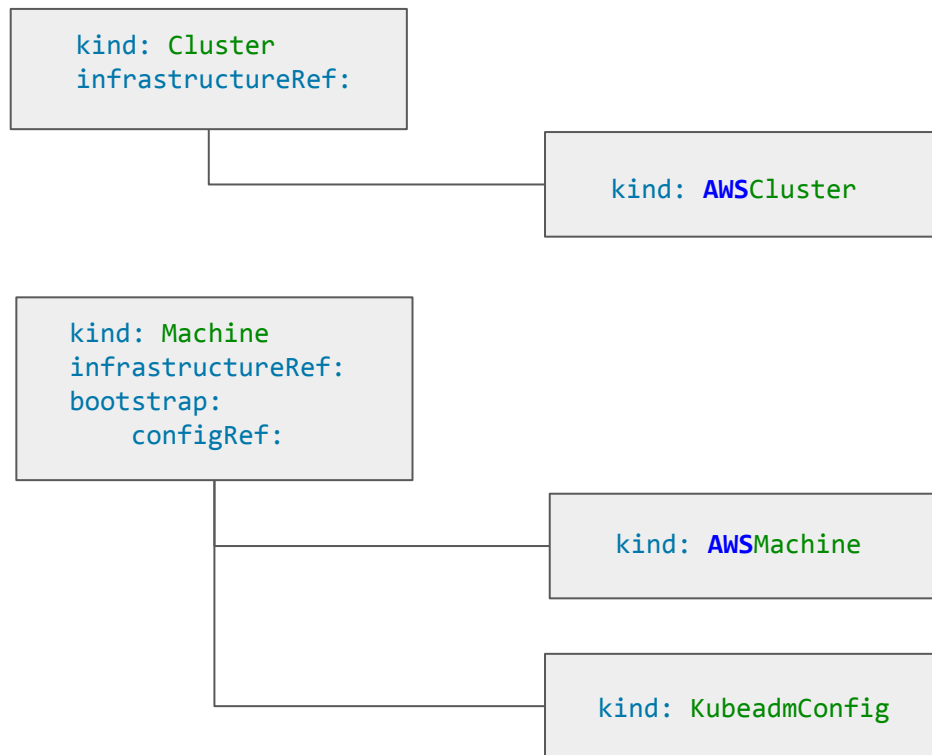
## node bootstrapper for my-machine

```
apiVersion:
bootstrap.cluster.x-k8s.io/v1alpha2
kind: KubeadmConfig
metadata:
  name: my-machine-bootstrapper
spec:
  initConfiguration: ...
  clusterConfiguration: ...
```



kubernetes

# Recap



# Trivia Time!

## Kubernetes Architecture

1. Name three Kubernetes controllers?



# Controllers in Kubernetes



In Kubernetes **operational best practices** for managing API objects are encoded into the controllers e.g.

*The deployment controller encodes best practices for Pod rollout strategies*



# Controllers in Cluster API



In cluster API **operational best practices** for *managing*\* a cluster are encoded into a set of controllers.

\*both provisioning and day 2 operations



kubernetes

# In practice: Cluster controllers



kind: Cluster

ClusterController

Reconcile cluster state  
(cluster infrastructure first, then Machines)

kind: AWSCluster

AWSClusterController

Reconcile cluster infrastructure on AWS  
(VPC, Elastic IP, ELB ...)

kind: VSphereCluster

VSphereClusterController

Reconcile cluster infrastructure on vSphere

kind: AzureCluster

AzureClusterController

Reconcile cluster infrastructure on Azure

...

...



kubernetes

# In practice: Machine Controllers



`kind: Machine`

MachineController

Reconcile machine state  
(bootstrap first, then infrastructure)

----- `kind: AWSMachine`  
----- ...

----- AWSMachineController  
----- ...

Reconcile machine infrastructures on AWS  
(EC2)

----- `kind: kubeadmConfig`

----- kubeadmConfigController

Generates the cloud-init scripts for  
bootstrapping the node with kubeadm

# Recap



One controller for each “Cluster” kind

One controller for each “Machine” kind

One controller for each “Config” kind



# Too many moving parts?



Up to a certain extent, yes, but

- as a user, you usually not care about controllers
- as an administrator
  - you can install only the providers you care about
  - controllers are packed into binaries/containers
- from an architectural PoV
  - each controller does one task (linux philosophy)
  - the overall system can scale
  - the overall system is more resilient



kubernetes

One last thing to know  
about how Cluster API  
works



# Mutable infrastructure



In a traditional infrastructure, servers are continually updated and modified in place.

Administrators SSH into their servers and apply changes on a server-by-server basis

# Immutable infrastructure



Immutable infrastructure is an approach to managing services and software deployments on IT resources wherein components are replaced rather than changed

# What about Cluster API approach?



Also in this case, Cluster API mirrors what Kubernetes does for Pods/Containers.

Machines (and the related K8s) are assumed immutable and replaced in case of necessity.




kubernetes

TL;DR;

Clusters, Machines, are only the beginning of the story ...

A new wave of tools and higher level abstractions is going to be created on top of that. *Stay tuned!*



# Agenda

- Intro
- How it works
- Demo




<https://cluster-api.sigs.k8s.io/user/quick-start.html>





# The chicken and egg problem



**wait, wait,  
I need a cluster for creating a new cluster...**

# 1. Create a cluster

```
# create a kind config file with a mount required by the docker infrastructure provider
cat > kind-cluster-with-extramounts.yaml <<EOF
kind: Cluster
apiVersion: kind.sigs.k8s.io/v1alpha3
nodes:
- role: control-plane
  extraMounts:
    - hostPath: /var/run/docker.sock
      containerPath: /var/run/docker.sock
EOF

# create a Kubernetes cluster
kind create cluster --config ./kind-cluster-with-extramounts.yaml

# makes the cluster accessible to kubectl
export KUBECONFIG="$(kind get kubeconfig-path)"
```



## 2. Transform the cluster into a management cluster.

```
# install cluster API provider
kubectl create -f
https://github.com/kubernetes-sigs/cluster-api/releases/download/v0.2.6/cluster-api-components.yaml

# install the kubeadm bootstrap provider
...

# install the docker infrastructure provider (for local testing)
...

# install the AWS infrastructure provider
...
```

or ... in the next release of cluster API


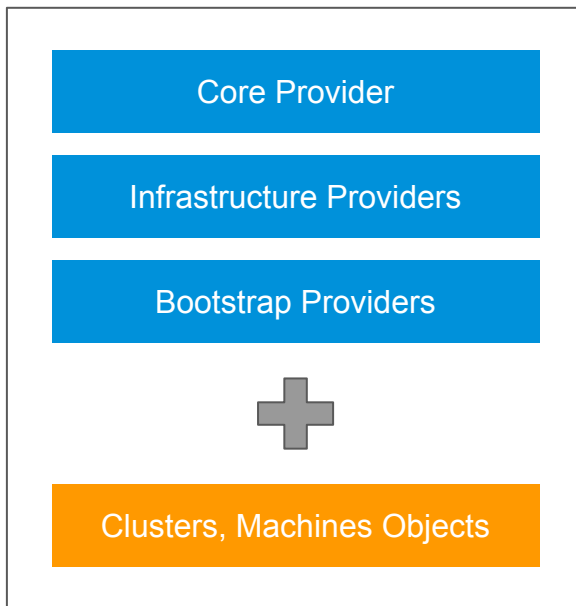
```
clusterctl init --infrastructure aws,docker --bootstrap kubeadm
```




# Wrapping up, so far ... and moving forward



## 1 Management cluster



## n Workload clusters



kubernetes

# Let's create our first workload cluster in docker !

```
# define the specs for the cluster
```

```
cat > cluster.yaml <<EOF
```

```
...
```

```
EOF
```

```
cat > controlplane.yaml <<EOF
```

```
...
```

```
EOF
```

```
cat > workers.yaml <<EOF
```

```
...
```

```
EOF
```

```
# create the cluster
```

```
kubectl apply -f cluster.yaml
```

```
kubectl apply -f controlplane.yaml
```

```
kubectl apply -f workers.yaml
```

or ... in the next release of cluster API

```
clusterctl config cluster my-cluster1 --infrastructure docker | kubectl apply -f -
```



# Let's create our first workload cluster in AWS !

```
# Guess what... you already know how to do it !
```



# Thank you!

Next up...  
implementing a Kubernetes controller

